

ApplianceTelemetryCorrelationAnalysis

October 20, 2024

0.1 Provided as-is (w/o support)

Kubernetes clusters collect various application and infrastructure statistics. While this information is useful, it's very difficult to identify which metrics are useful for monitoring and troubleshooting. The Goal here is to collect this information, and use a statistical model to identify which metrics should be included in reports/dashboard such that: * Unnecessary overhead and sensory overload can be reduced. * Time can be saved by prioritising monitoring the correct metrics.

This process needs to assume **zero knowledge** of the workings of the cluster, workload being run and any other information. This way, **generic** clusters can be monitored without explicitly programming dashboards based on internal knowledge. This is also a good method to discover/verify application knowledge/bottlenecks with statistical data analysis.

0.2 Step1: Data Loading

We will load cpu, memory, task_queue information along with stats from structured and unstructured scans from csv files stored on disk using the `dataframeLoader` helper.

*# The dataframeLoader helper function implements the loadApplianceTimeSeriesData method.
This method loads the csv files, and pivots them to generate distinct "metrics" timeseries.
see <https://github.com/amitgupta7/docker-jupy-ntbk-s3-reporting/blob/main/dataframeLoader.py>*

```
[1]: import sys
      sys.path.append('../')

      import dataframeLoader as dfl
      import pandas as pd
      from importlib import reload
      reload(dfl)

      # Provide csv data location and appliance and timerange information.
      root = '../..dataDir'
      fromDt = '2024-06-01'
      toDt = '2024-11-01'

      # Provide list of prometheus metrics to load.
      metricsArr = ['cpu_used', 'download_workers_count', 'memory_used',
                    ↪ 'task_queue_length', 'infra_access_latency', 'pod_cpu_usage',
                    ↪ 'pod_memory_usage']
      # metricsArr = ['cpu_used'
```

```

#             , 'task_queue_length'
#             , 'memory_used'
#             ]

daterange=[fromDt, toDt]
df = dfl.loadApplianceTimeSeriesData(root, metricsArr, daterange)

```

```

loading Unstrctured Data from file: SCANPROC-*.csv
loading Strctured Data from file: STRUCTURED-*.csv
processing securiti_appliance_cpu_used-max*.csv
processing securiti_appliance_cpu_used-avg*.csv
processing securiti_appliance_download_workers_count-max*.csv
processing securiti_appliance_download_workers_count-avg*.csv
processing securiti_appliance_memory_used-max*.csv
processing securiti_appliance_memory_used-avg*.csv
processing securiti_appliance_task_queue_length-max*.csv
processing securiti_appliance_task_queue_length-avg*.csv
processing securiti_appliance_infra_access_latency-max*.csv
processing securiti_appliance_infra_access_latency-avg*.csv
processing securiti_appliance_pod_cpu_usage-max*.csv
processing securiti_appliance_pod_cpu_usage-avg*.csv
processing securiti_appliance_pod_memory_usage-max*.csv
processing securiti_appliance_pod_memory_usage-avg*.csv
loading Unstrctured Data from file: UNSTRUCTURED-*.csv

```

0.3 Step2: Data Pivoting

We now aggregate the data by `appliance_id` (unique identifier for our cluster) and `ts` timestamp, to get different metrics values as separate columns. Notice there are: * Statistical significance (consider correlation for appliances with atleast 30 non-zero scan time values, ie the appliance should be scanning for atleast 30 hours before considering it statistically significant). * 200+ appliances * Total scanning time of over 15 years!!

- 33 metrics -> 22 metrics
 - Decide between `max` or `avg` values if both are present.
 - We chose to display `avg` values metrics in this case after some trial and error.
 - * Except for memory (where `max` indicates spikes/oom conditions better)
- Tracked every hour

```
[2]: #consider only appliances with a certain number of scanning intervals.
min_scanning_intervals = 10
dfp = df.pivot_table(index=['appliance_id','ts'], columns=['metrics'],
    ↪values='value', aggfunc='sum').reset_index()
max_list = list(dfp.filter(regex='max'))
# maxlist = ['cpu_used_max', 'linkerq_max', 'memory_used_avg', 'taskq_max',
    ↪'tmp_taskq_max']
# we would like to use max memory (to indicate spikes/oom conditions)
max_list = list(map(lambda x: x.replace('memory_used_max', 'memory_used_avg'),
    ↪max_list))
print('cols removed', max_list)
dfp = dfp[dfp.columns.drop(max_list)]
stat_sig = dfp.fillna(0).groupby(['appliance_id']).scanTime.agg(lambda x: x.
    ↪ne(0).sum())
stat_sig = stat_sig[stat_sig > min_scanning_intervals].
    ↪sort_values(ascending=False)
print('cols retained', dfp.columns)
print(len(stat_sig), 'statistically significant appliances with total scan time
    ↪of', stat_sig.sum()/24/365, 'years')
dfp = dfp[dfp.appliance_id.isin(stat_sig.index)]

display(dfp)
```

```
cols removed ['cpu_used_max', 'download_workers_count_max', 'esLatency_max',
'linkerq_max', 'memory_used_avg', 'pgLatency_max', 'pod_cpu_usage_max',
'pod_memory_usage_max', 'redisLatency_max', 'taskq_max', 'tmp_taskq_max']
cols retained Index(['appliance_id', 'ts', 'IdleTimeInHrs', 'avgFileSizeInMB',
'cpu_used_avg', 'dataScannedinGB', 'download_workers_count_avg',
'esLatency_avg', 'fileDownloadTimeInHrs', 'linkerq_avg',
'memory_used_max', 'numFilesScanned', 'numberOfChunksScanned',
'numberOfColsScanned', 'pgLatency_avg', 'pod_cpu_usage_avg',
'pod_memory_usage_avg', 'redisLatency_avg', 'scanTime', 'taskq_avg',
'tmp_taskq_avg', 'uniqPodCount'],
      dtype='object', name='metrics')
```

277 statistically significant appliances with total scan time of
15.499086757990868 years

metrics	appliance_id	ts \
57	01c75278-9c0d-41be-b693-c970b18dbedc	2024-06-01 00:00:00
58	01c75278-9c0d-41be-b693-c970b18dbedc	2024-06-01 01:00:00
59	01c75278-9c0d-41be-b693-c970b18dbedc	2024-06-01 02:00:00
60	01c75278-9c0d-41be-b693-c970b18dbedc	2024-06-01 03:00:00
61	01c75278-9c0d-41be-b693-c970b18dbedc	2024-06-01 04:00:00
...
848488	ffc2fb2f-52e0-42f3-8564-9545dbc6b747	2024-08-21 00:00:00
848489	ffc2fb2f-52e0-42f3-8564-9545dbc6b747	2024-08-21 01:00:00
848490	ffc2fb2f-52e0-42f3-8564-9545dbc6b747	2024-08-21 02:00:00
848491	ffc2fb2f-52e0-42f3-8564-9545dbc6b747	2024-08-21 03:00:00

848492 ffc2fb2f-52e0-42f3-8564-9545dbc6b747 2024-08-21 04:00:00

metrics	IdleTimeInHrs	avgFileSizeInMB	cpu_used_avg	dataScannedinGB	\
57	2.396345	0.036537	NaN	0.001096	
58	1.804338	0.851661	NaN	0.034066	
59	5.472733	0.067678	NaN	0.006091	
60	2.701624	6.814020	NaN	0.293003	
61	3.396194	0.304974	NaN	0.015249	
...	
848488	NaN	NaN	2.723333	NaN	
848489	NaN	NaN	2.723333	NaN	
848490	NaN	NaN	2.723333	NaN	
848491	NaN	NaN	2.723333	NaN	
848492	NaN	NaN	2.723333	NaN	

metrics	download_workers_count_avg	esLatency_avg	fileDownloadTimeInHrs	\
57	NaN	NaN	NaN	
58	NaN	NaN	NaN	
59	NaN	NaN	NaN	
60	NaN	NaN	NaN	
61	NaN	NaN	NaN	
...	
848488	NaN	0.000603	NaN	
848489	NaN	0.000603	NaN	
848490	NaN	0.000603	NaN	
848491	NaN	0.000603	NaN	
848492	NaN	0.000603	NaN	

metrics	linkerq_avg	...	numberOfChunksScanned	numberOfColsScanned	\
57	NaN	...	NaN	NaN	
58	NaN	...	NaN	NaN	
59	NaN	...	NaN	NaN	
60	NaN	...	NaN	NaN	
61	NaN	...	NaN	NaN	
...	
848488	NaN	...	NaN	NaN	
848489	NaN	...	NaN	NaN	
848490	NaN	...	NaN	NaN	
848491	NaN	...	NaN	NaN	
848492	NaN	...	NaN	NaN	

metrics	pgLatency_avg	pod_cpu_usage_avg	pod_memory_usage_avg	\
57	NaN	NaN	NaN	
58	NaN	NaN	NaN	
59	NaN	NaN	NaN	
60	NaN	NaN	NaN	
61	NaN	NaN	NaN	
...	

848488	0.056978	4.813333	36.355
848489	0.056978	4.813333	36.355
848490	0.056978	4.813333	36.355
848491	0.056978	4.813333	36.355
848492	0.056978	4.813333	36.355

metrics	redisLatency_avg	scanTime	taskq_avg	tmp_taskq_avg	uniqPodCount
57	NaN	0.029513	NaN	NaN	3.0
58	NaN	0.195418	NaN	NaN	3.0
59	NaN	0.057962	NaN	NaN	4.0
60	NaN	0.196190	NaN	NaN	3.0
61	NaN	0.035388	NaN	NaN	3.0
...
848488	0.003052	NaN	NaN	NaN	NaN
848489	0.003052	NaN	NaN	NaN	NaN
848490	0.003052	NaN	NaN	NaN	NaN
848491	0.003052	NaN	NaN	NaN	NaN
848492	0.003052	NaN	NaN	NaN	NaN

[470148 rows x 22 columns]

0.4 Step 3: Data transformation and correlation

We need to achieve two main goals: 1. Isolate data for individual appliance. 2. Remove ghost correlation between unrelated metrics. * We will calculate percentage change between adjacent timeseries values. 3. Calculate absolute correlation between metrics for each single appliance. * Transpose every metrics correlation. 4. Generate correlation for every `appliance_id` and `metric` identifier using steps 1, 2 and 3

```
[3]: # appliance = '01c75278-9c0d-41be-b693-c970b18dbedc'
# for metric in metrics_category_order:
dfc_arr = []
for pod in dfp.appliance_id.unique():
    dfa = dfp[(dfp.appliance_id == pod)]
    dfa = dfa.drop(['appliance_id', 'ts'], axis=1)
    dfa = dfa.pct_change(periods=1, fill_method=None)
    dfca = dfa.corr().abs()
    # print(type(dfca))
    for col in dfca.columns:
        # print(col)
        dfc = dfca[col].to_frame().T
        dfc.insert(0, 'metric', col)
        dfc.insert(0, 'appliance_id', pod)
        dfc_arr.append(dfc)
dfc = pd.concat(dfc_arr, ignore_index=True)
dfc.set_index('appliance_id', inplace=True)
dfc.head()
```

```

[3]: metrics
appliance_id
01c75278-9c0d-41be-b693-c970b18dbedc IdleTimeInHrs
01c75278-9c0d-41be-b693-c970b18dbedc avgFileSizeInMB
01c75278-9c0d-41be-b693-c970b18dbedc cpu_used_avg
01c75278-9c0d-41be-b693-c970b18dbedc dataScannedinGB
01c75278-9c0d-41be-b693-c970b18dbedc download_workers_count_avg

metrics IdleTimeInHrs avgFileSizeInMB \
appliance_id
01c75278-9c0d-41be-b693-c970b18dbedc 1.000000 0.032743
01c75278-9c0d-41be-b693-c970b18dbedc 0.032743 1.000000
01c75278-9c0d-41be-b693-c970b18dbedc 0.022182 0.093800
01c75278-9c0d-41be-b693-c970b18dbedc 0.134145 0.558588
01c75278-9c0d-41be-b693-c970b18dbedc NaN NaN

metrics cpu_used_avg dataScannedinGB \
appliance_id
01c75278-9c0d-41be-b693-c970b18dbedc 0.022182 0.134145
01c75278-9c0d-41be-b693-c970b18dbedc 0.093800 0.558588
01c75278-9c0d-41be-b693-c970b18dbedc 1.000000 0.167504
01c75278-9c0d-41be-b693-c970b18dbedc 0.167504 1.000000
01c75278-9c0d-41be-b693-c970b18dbedc NaN NaN

metrics download_workers_count_avg \
appliance_id
01c75278-9c0d-41be-b693-c970b18dbedc NaN
01c75278-9c0d-41be-b693-c970b18dbedc NaN
01c75278-9c0d-41be-b693-c970b18dbedc NaN
01c75278-9c0d-41be-b693-c970b18dbedc NaN
01c75278-9c0d-41be-b693-c970b18dbedc NaN

metrics esLatency_avg fileDownloadTimeInHrs \
appliance_id
01c75278-9c0d-41be-b693-c970b18dbedc 0.015515 0.337686
01c75278-9c0d-41be-b693-c970b18dbedc 0.028873 0.063530
01c75278-9c0d-41be-b693-c970b18dbedc 0.000113 0.010405
01c75278-9c0d-41be-b693-c970b18dbedc 0.007807 0.037724
01c75278-9c0d-41be-b693-c970b18dbedc NaN NaN

metrics linkerq_avg memory_used_max ... \
appliance_id
01c75278-9c0d-41be-b693-c970b18dbedc NaN 0.323574 ...
01c75278-9c0d-41be-b693-c970b18dbedc NaN 0.362822 ...
01c75278-9c0d-41be-b693-c970b18dbedc NaN 0.596260 ...
01c75278-9c0d-41be-b693-c970b18dbedc NaN 0.598367 ...
01c75278-9c0d-41be-b693-c970b18dbedc NaN NaN ...

```

metrics	numberOfChunksScanned	\
appliance_id		
01c75278-9c0d-41be-b693-c970b18dbedc	0.012446	
01c75278-9c0d-41be-b693-c970b18dbedc	0.028560	
01c75278-9c0d-41be-b693-c970b18dbedc	0.094911	
01c75278-9c0d-41be-b693-c970b18dbedc	0.155238	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	

metrics	numberOfColsScanned	pgLatency_avg	\
appliance_id			
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	0.066219	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	0.059433	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	0.022957	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	0.038979	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	NaN	

metrics	pod_cpu_usage_avg	pod_memory_usage_avg	\
appliance_id			
01c75278-9c0d-41be-b693-c970b18dbedc	0.010472	0.007449	
01c75278-9c0d-41be-b693-c970b18dbedc	0.048583	0.032665	
01c75278-9c0d-41be-b693-c970b18dbedc	0.013278	0.022344	
01c75278-9c0d-41be-b693-c970b18dbedc	0.057678	0.024285	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	NaN	

metrics	redisLatency_avg	scanTime	taskq_avg	\
appliance_id				
01c75278-9c0d-41be-b693-c970b18dbedc	0.002635	0.005601	0.009926	
01c75278-9c0d-41be-b693-c970b18dbedc	0.007076	0.094383	0.052507	
01c75278-9c0d-41be-b693-c970b18dbedc	0.020036	0.017485	0.874397	
01c75278-9c0d-41be-b693-c970b18dbedc	0.047678	0.231155	0.028069	
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	NaN	NaN	

metrics	tmp_taskq_avg	uniqPodCount
appliance_id		
01c75278-9c0d-41be-b693-c970b18dbedc	0.012718	0.052765
01c75278-9c0d-41be-b693-c970b18dbedc	0.020660	0.007626
01c75278-9c0d-41be-b693-c970b18dbedc	0.302852	0.023543
01c75278-9c0d-41be-b693-c970b18dbedc	0.022780	0.582436
01c75278-9c0d-41be-b693-c970b18dbedc	NaN	NaN

[5 rows x 21 columns]

0.5 Step 4: Isolate related metrics using correlation

We now iterate over each `metric`, to see if there is any significant statistical correlation to be found across `appliance_ids`. This is done with two steps:

1. Removing outliers:
 - Remove any metrics with 3rd quantile correlation value below the cut-off. This cut-off can be varied for depending on use cases:
 - 0.9 for Exec Dashboards
 - 0.8 for Customer Ops
 - 0.7 for L1 - support
 - 0.6 for L2 - suport

Please note that we are filtering metrics with 3rd quantile correlation below the moderate cut-off. This ensures that atleast 25% of the values are correlated to reduce outliers.

2. Plot box chart to visually represent metrics with any correlation (for cutoff as 0.3).
3. The network graph indicates specific correlation edges between metrics.

0.6 Final List of metrics

The below table shows the list of **metrics** that are useful with respective correlation **cutoff**. The cut-off values can be interpreted as follows:

- below 0.3 negligible correlation
- 0.3 to 0.5 Low positive (negative) correlation
- 0.5 to 0.7 Moderate positive (negative) correlation
- 0.7 to 0.9 High positive (negative) correlation
- 0.9 to 0.1 Very High positive (negative) correlation

```
[4]: import gravis as gv
import itertools
import networkx as nx
from IPython.display import Image

corr_vals = [0.6, 0.7, 0.8, 0.9]
mtrx_arr = []
graph_arr = []
for cutoff in corr_vals:
    arr = []
    for metr in dfc.metric.unique():
        dfcm = dfc[(dfc.metric == metr)]
        dfcm = dfcm.drop('metric', axis=1)
        dfcm = dfcm.drop(metr, axis=1)
        dfcm = dfcm.dropna(axis = 0, how = 'all')
        dfcm = dfcm.loc[:, dfcm.quantile(q=0.75) > cutoff]
        for x in dfcm.columns:
            arr.append(x)
            graph_arr.append((metr, x))
    if(cutoff == corr_vals[0]):
        if len(dfcm.columns) > 0:
            title=f'''Absolute correlation vs percent-change of {metr}
            (For median correlation greater than {cutoff})
            '''
```



```

        dfcm.plot(kind='box',vert=False,title=title,colormap='tab20')
arr = list(set(arr))
arr.insert(0,cutoff)
mtrx_arr.append(arr)

display(pd.DataFrame(list(map(list, itertools.zip_longest(*mtrx_arr,
↪fillvalue="")))))

g = nx.DiGraph()
g.add_edges_from(graph_arr)
fig = gv.d3(g
    , graph_height=800
    , use_node_size_normalization=True
    , zoom_factor=2
    , node_size_normalization_max=30
    ,use_edge_size_normalization=True
    ,use_collision_force=True
    ,node_label_size_factor=0.8
    , layout_algorithm_active=True
    )
fig.export_jpg('graph2.jpg', overwrite=True)
print("Correlation graph between appliance metrics")
Image('graph2.jpg')
# fig.display()

```

	0	1 \
0	0.6	0.7
1	avgFileSizeInMB	avgFileSizeInMB
2	memory_used_max	memory_used_max
3	scanTime	scanTime
4	uniqPodCount	dataScannedinGB
5	dataScannedinGB	IdleTimeInHrs
6	IdleTimeInHrs	download_workers_count_avg
7	download_workers_count_avg	numberOfColsScanned
8	numberOfColsScanned	cpu_used_avg
9	cpu_used_avg	taskq_avg
10	taskq_avg	numberOfChunksScanned
11	numberOfChunksScanned	fileDownloadTimeInHrs
12	fileDownloadTimeInHrs	linkerq_avg
13	linkerq_avg	numFilesScanned
14	numFilesScanned	

	2	3
0	0.8	0.9
1	avgFileSizeInMB	scanTime
2	scanTime	numberOfColsScanned
3	dataScannedinGB	numberOfChunksScanned
4	download_workers_count_avg	fileDownloadTimeInHrs

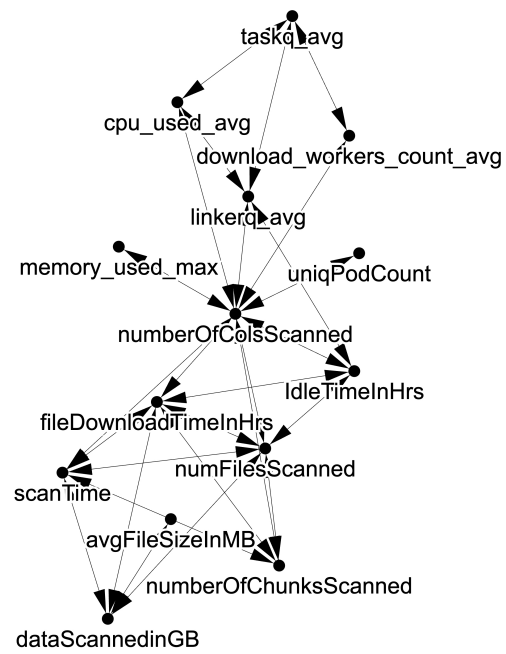
```

5         numberOfColsScanned        numFilesScanned
6     numberOfChunksScanned
7     fileDownloadTimeInHrs
8         numFilesScanned
9
10
11
12
13
14

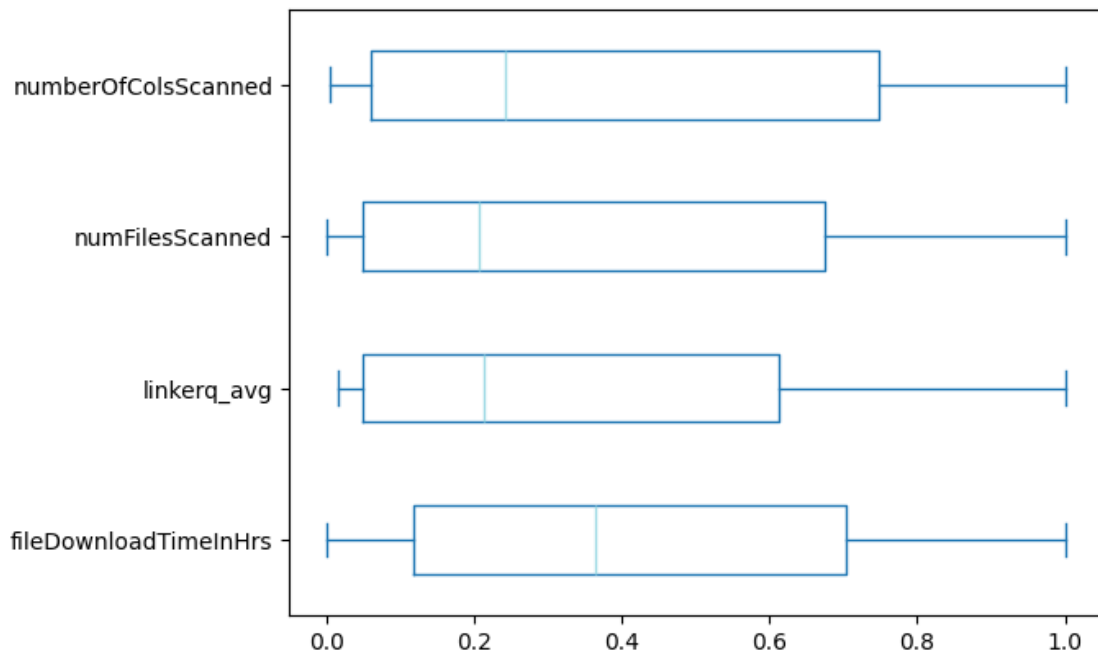
```

Correlation graph between appliance metrics

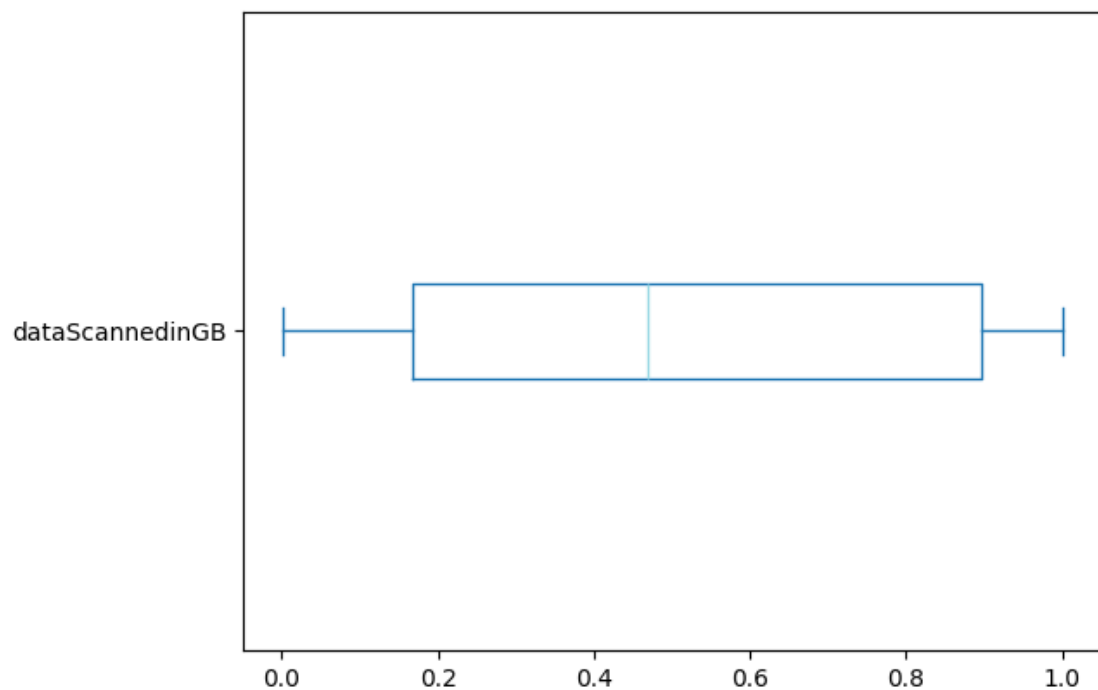
[4]:



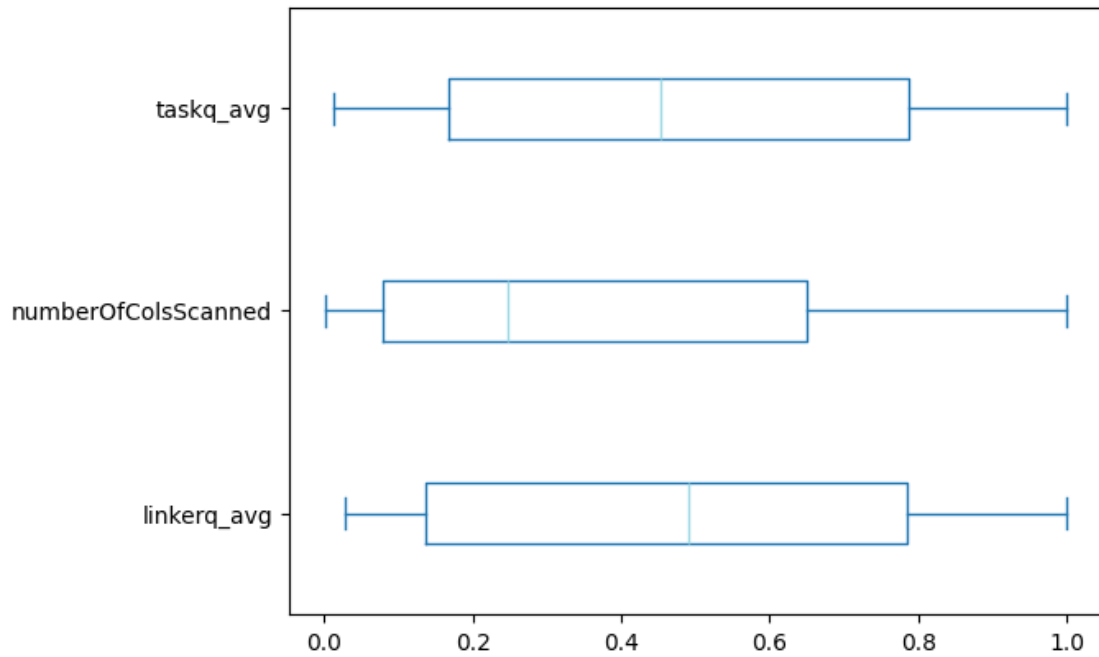
Absolute correlation vs percent-change of IdleTimeInHrs
(For median correlation greater than 0.6)



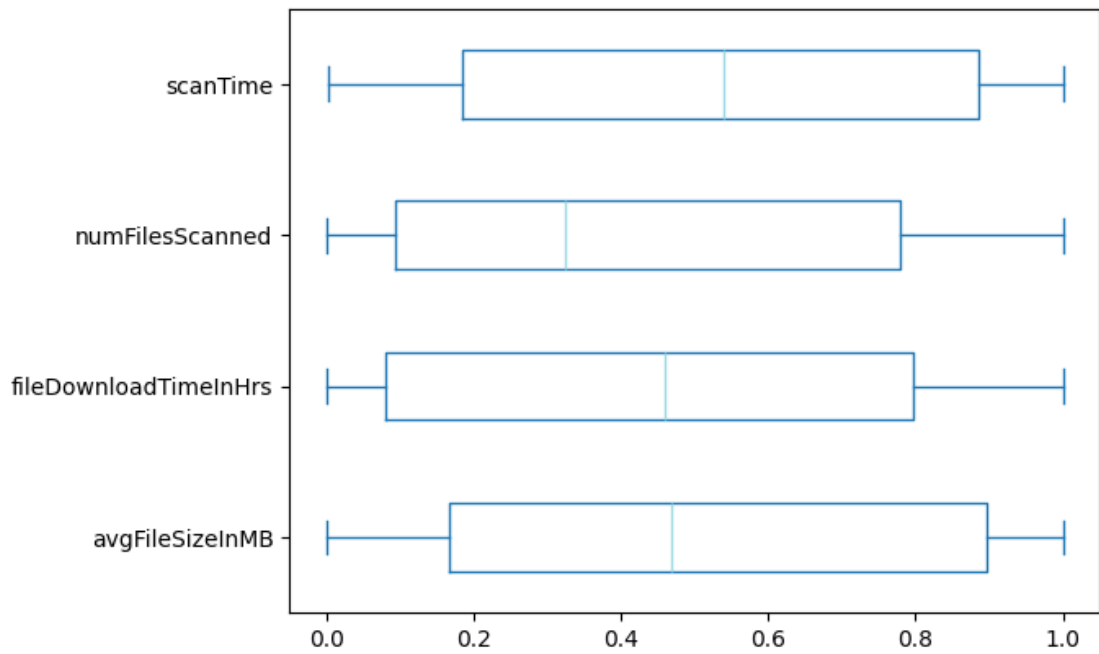
Absolute correlation vs percent-change of avgFileSizeInMB
(For median correlation greater than 0.6)



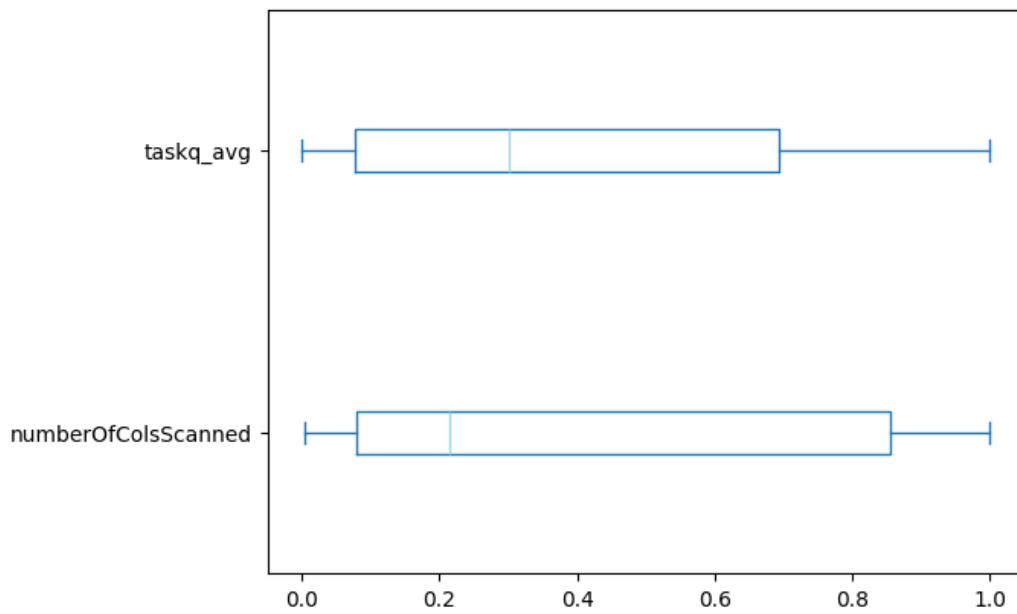
Absolute correlation vs percent-change of `cpu_used_avg`
(For median correlation greater than 0.6)



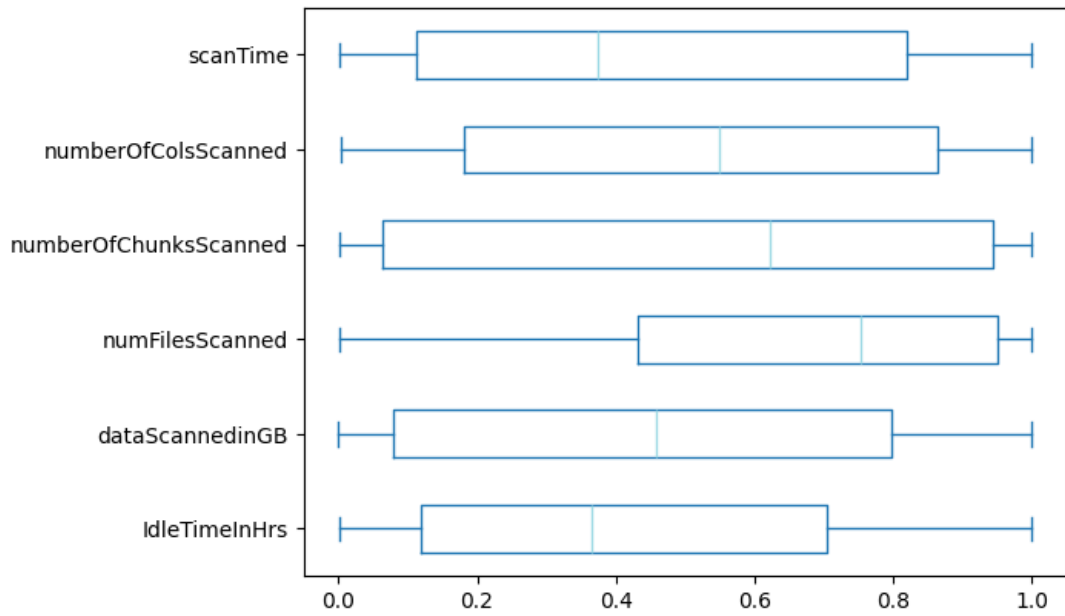
Absolute correlation vs percent-change of dataScannedinGB
(For median correlation greater than 0.6)



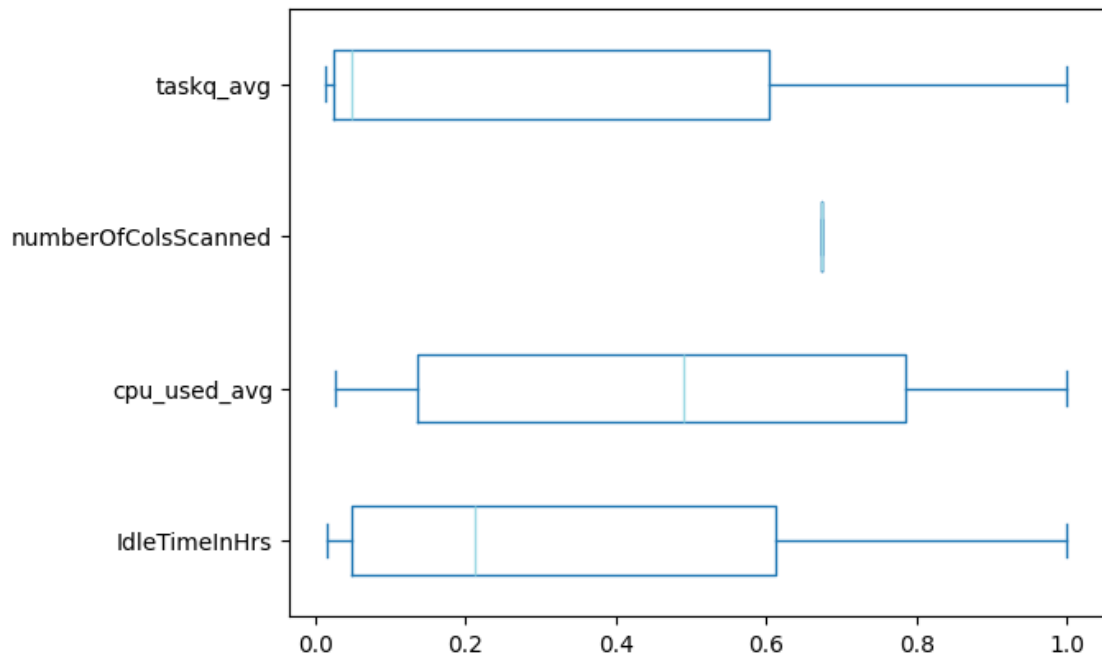
Absolute correlation vs percent-change of download_workers_count_avg
(For median correlation greater than 0.6)



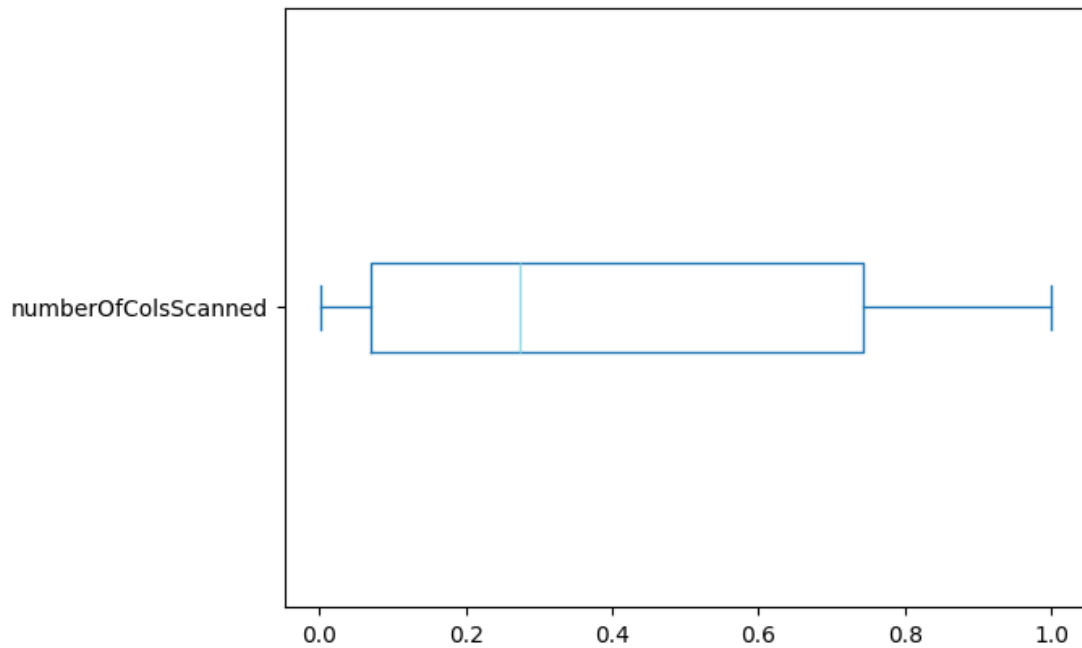
Absolute correlation vs percent-change of fileDownloadTimeInHrs
(For median correlation greater than 0.6)



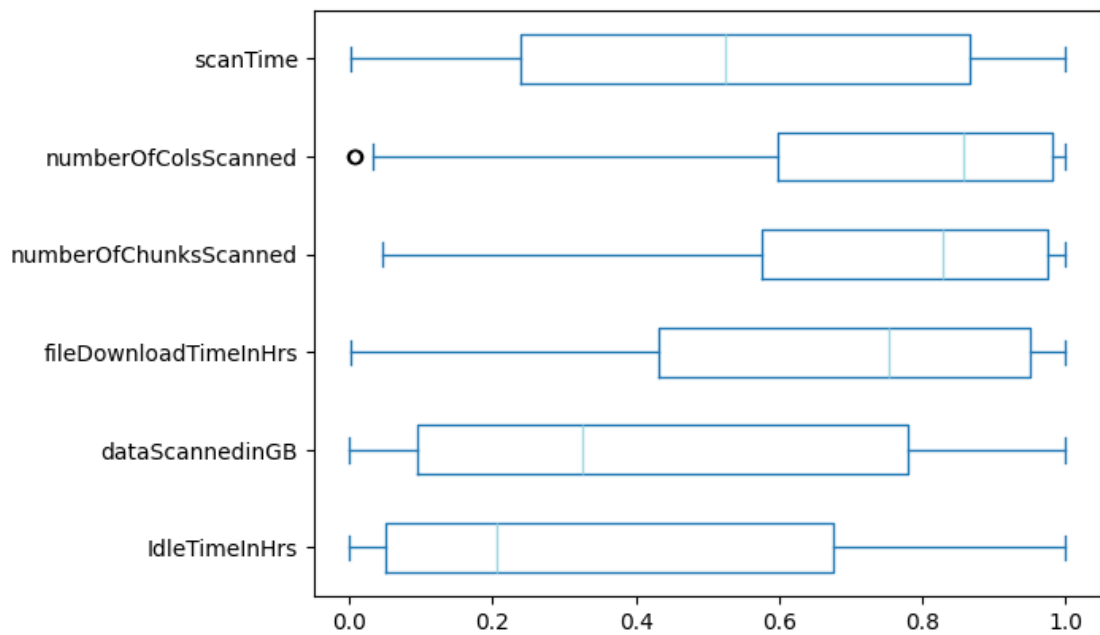
Absolute correlation vs percent-change of linkerq_avg
(For median correlation greater than 0.6)



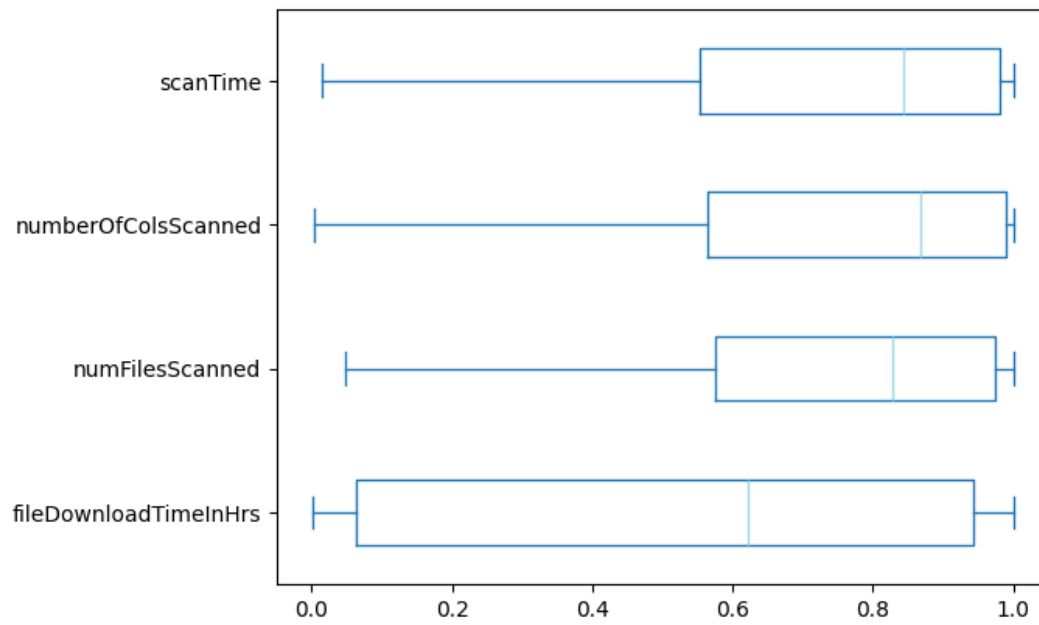
Absolute correlation vs percent-change of memory_used_max
(For median correlation greater than 0.6)



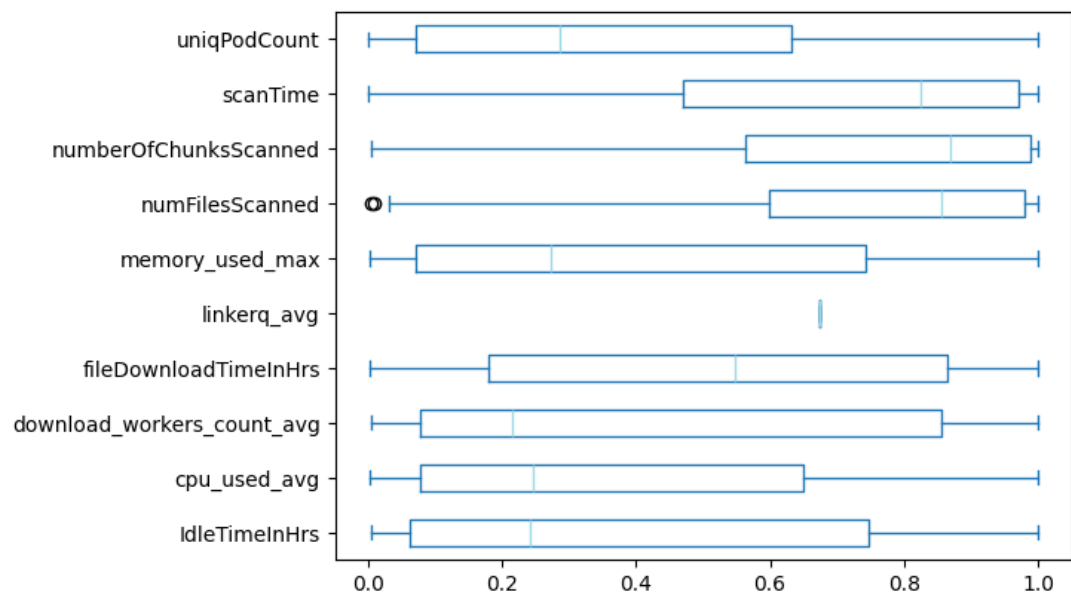
Absolute correlation vs percent-change of numFilesScanned
(For median correlation greater than 0.6)



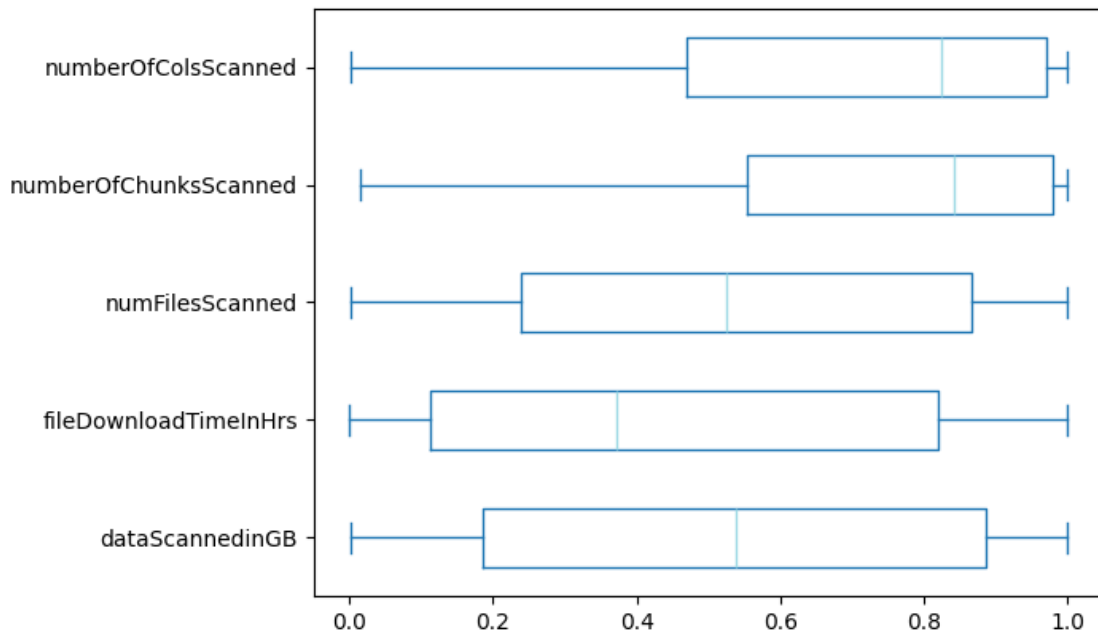
Absolute correlation vs percent-change of numberOfChunksScanned
(For median correlation greater than 0.6)



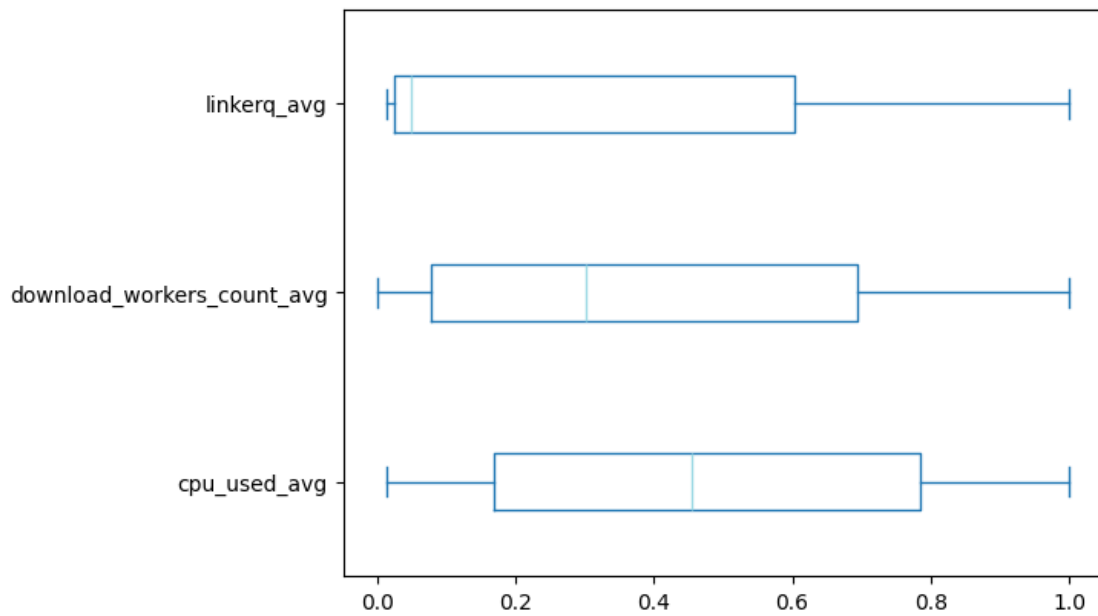
Absolute correlation vs percent-change of numberOfColsScanned
(For median correlation greater than 0.6)



Absolute correlation vs percent-change of scanTime
(For median correlation greater than 0.6)



Absolute correlation vs percent-change of taskq_avg
(For median correlation greater than 0.6)



Absolute correlation vs percent-change of uniqPodCount
(For median correlation greater than 0.6)

